

## Homework 2 Report

Name: Viprav Lipare

Student ID: 801288922

Homework Number: 2

Github Repository: <https://github.com/vipravlipare/ECGR-5105-Intro-to-Machine-Learning>

### Problem 1A

#### Source Code (1A):

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import MinMaxScaler, StandardScaler
file_path = "Housing.csv"
sample = df = pd.read_csv(file_path)

inputs = ['area', 'bedrooms', 'bathrooms', 'stories', 'parking'];
X = sample[inputs].values
Y = sample.values[:, 0]; #Price

m = len(Y)
X_0 = np.ones((m, 1))
X = np.hstack((X_0, X))
theta = np.zeros(X.shape[1])
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.2,
random_state=42)

# defining function for computing the cost for linear regression
def compute_cost(X, y, theta):
    m = len(y)
    predictions = X.dot(theta)
    errors = predictions - y
    sqrErrors = np.square(errors)
    J = 1 / (2 * m) * np.sum(sqrErrors)
    return J

# defining function for gradient descent algorithm
def gradient_descent(Xt, Yt, Xv, Yv, theta, alpha, iterations):
    train_cost_history = np.zeros(iterations)
    valid_cost_history = np.zeros(iterations)
    m = len(Yt)
    for i in range(iterations):
```

```

predictions = Xt.dot(theta)
errors = np.subtract(predictions, Yt)
sum_delta = (alpha / m) * Xt.transpose().dot(errors)
theta = theta - sum_delta
train_cost_history[i] = compute_cost(Xt, Yt, theta)
valid_cost_history[i] = compute_cost(Xv, Yv, theta)
return theta, train_cost_history, valid_cost_history

theta = np.zeros(X.shape[1])
iterations = 1000

alpha = 0.1
theta_1, train_cost_history_1, valid_cost_history_1 =
gradient_descent(X_train, Y_train, X_test, Y_test, theta, alpha,
iterations) # 0.1 Learning Rate

alpha = 0.01
theta_2, train_cost_history_2, valid_cost_history_2 =
gradient_descent(X_train, Y_train, X_test, Y_test, theta, alpha,
iterations) # 0.01 Learning Rate

plt.plot(range(1, iterations+1), train_cost_history_1, color='blue',
label='Training Loss (0.1)')
plt.plot(range(1, iterations+1), valid_cost_history_1, color='red',
label='Validation Loss (0.1)')
plt.plot(range(1, iterations+1), train_cost_history_2, color='blue',
linestyle='dotted', label='Training Loss (0.01)')
plt.plot(range(1, iterations+1), valid_cost_history_2, color='red',
linestyle='dotted', label='Validation Loss (0.01)')
plt.rcParams["figure.figsize"] = (10,6)
plt.grid()
plt.xlabel('Number of iterations')
plt.ylabel('Cost (J)')
plt.title('Training vs Validation Loss')
plt.legend()
plt.show()

features = ['X_0', 'area', 'bedrooms', 'bathrooms', 'stories', 'parking']


```

```

print("\nBest Model Parameters (Theta) for Problem 1A: Learning Rate = 0.1\n")
theta_list = theta_1.flatten()
for label, value in zip(features, theta_list):
    print(f"{label:>12}: {value:.2f}")

```

### Figures / Results (1A):



### **Problem 1B**

#### Source Code (1B):

```

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import MinMaxScaler, StandardScaler
file_path = "Housing.csv"
sample = df = pd.read_csv(file_path)

```

```

inputs = ['area', 'bedrooms', 'bathrooms', 'stories', 'mainroad',
'guestroom', 'basement', 'hotwaterheating', 'airconditioning', 'parking',
'prefarea']

X = sample[inputs].values
Y = sample.values[:, 0]; #Price

binary_cols = ['mainroad', 'guestroom', 'basement', 'hotwaterheating',
'airconditioning', 'prefarea']
for col in binary_cols:
    sample[col] = sample[col].map({'yes': 1, 'no': 0})

m = len(Y)
X_0 = np.ones((m, 1))
X_cols = [sample['area'].values.reshape(m, 1),
           sample['bedrooms'].values.reshape(m, 1),
           sample['bathrooms'].values.reshape(m, 1),
           sample['stories'].values.reshape(m, 1),
           sample['mainroad'].values.reshape(m, 1),
           sample['guestroom'].values.reshape(m, 1),
           sample['basement'].values.reshape(m, 1),
           sample['hotwaterheating'].values.reshape(m, 1),
           sample['airconditioning'].values.reshape(m, 1),
           sample['parking'].values.reshape(m, 1),
           sample['prefarea'].values.reshape(m, 1)]

X = np.hstack([X_0] + X_cols)

theta = np.zeros(X.shape[1])
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.2,
random_state=42)

# defining function for computing the cost for linear regression
def compute_cost(X, y, theta):
    m = len(y)
    predictions = X.dot(theta)
    errors = predictions - y
    sqrErrors = np.square(errors)
    J = 1 / (2 * m) * np.sum(sqrErrors)
    return J

# defining function for gradient descent algorithm

```

```

def gradient_descent(Xt, Yt, Xv, Yv, theta, alpha, iterations):
    train_cost_history = np.zeros(iterations)
    valid_cost_history = np.zeros(iterations)
    m = len(Yt)
    for i in range(iterations):
        predictions = Xt.dot(theta)
        errors = np.subtract(predictions, Yt)
        sum_delta = (alpha / m) * Xt.transpose().dot(errors)
        theta = theta - sum_delta
        train_cost_history[i] = compute_cost(Xt, Yt, theta)
        valid_cost_history[i] = compute_cost(Xv, Yv, theta)
    return theta, train_cost_history, valid_cost_history

theta = np.zeros(X.shape[1])
iterations = 1000

alpha = 0.1
theta_1, train_cost_history_1, valid_cost_history_1 =
gradient_descent(X_train, Y_train, X_test, Y_test, theta, alpha,
iterations) # 0.1 Learning Rate

alpha = 0.01
theta_2, train_cost_history_2, valid_cost_history_2 =
gradient_descent(X_train, Y_train, X_test, Y_test, theta, alpha,
iterations) # 0.01 Learning Rate

plt.plot(range(1, iterations+1), train_cost_history_1, color='blue',
label='Training Loss (0.1)')
plt.plot(range(1, iterations+1), valid_cost_history_1, color='red',
label='Validation Loss (0.1)')
plt.plot(range(1, iterations+1), train_cost_history_2, color='blue',
linestyle='dotted', label='Training Loss (0.01)')
plt.plot(range(1, iterations+1), valid_cost_history_2, color='red',
linestyle='dotted', label='Validation Loss (0.01)')
plt.rcParams["figure.figsize"] = (10,6)
plt.grid()
plt.xlabel('Number of iterations')
plt.ylabel('Cost (J)')
plt.title('Training vs Validation Loss')
plt.legend()

```

```


plt.show()

features = ['X_0', 'area', 'bedrooms', 'bathrooms', 'stories', 'mainroad',
'guestroom', 'basement', 'hotwaterheating', 'airconditioning', 'parking',
'prefarea']

print("\nBest Model Parameters (Theta) for Problem 1B: Learning Rate =
0.1\n")
theta_list = theta_1.flatten()
for label, value in zip(features, theta_list):
    print(f"{label:>12}: {value:.2f}")

```

### Figures / Results (1B):



### **Problem 2A**

### Source Code (2A):

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import MinMaxScaler, StandardScaler
file_path = "Housing.csv"
sample = df = pd.read_csv(file_path)
inputs = ['area', 'bedrooms', 'bathrooms', 'stories', 'parking'];

X = sample[inputs].values
Y = sample.values[:, 0]; #Price
m = len(Y)
X_0 = np.ones((m, 1))
X = np.hstack((X_0, X))
theta = np.zeros(X.shape[1])

scaler = MinMaxScaler()
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.2,
random_state=42)
X_train[:, 1:] = scaler.fit_transform(X_train[:, 1:])
X_test[:, 1:] = scaler.transform(X_test[:, 1:])

# defining function for computing the cost for linear regression
def compute_cost(X, y, theta):
    m = len(y)
    predictions = X.dot(theta)
    errors = predictions - y
    sqrErrors = np.square(errors)
    J = 1 / (2 * m) * np.sum(sqrErrors)
    return J

# defining function for gradient descent algorithm
def gradient_descent(Xt, Yt, Xv, Yv, theta, alpha, iterations):
    train_cost_history = np.zeros(iterations)
    valid_cost_history = np.zeros(iterations)
    m = len(Yt)
    for i in range(iterations):
        predictions = Xt.dot(theta)
        errors = np.subtract(predictions, Yt)
        theta = theta - ((alpha / m) * errors)
```

```

sum_delta = (alpha / m) * Xt.transpose().dot(errors)
theta = theta - sum_delta
train_cost_history[i] = compute_cost(Xt, Yt, theta)
valid_cost_history[i] = compute_cost(Xv, Yv, theta)
return theta, train_cost_history, valid_cost_history

theta = np.zeros(X.shape[1])
iterations = 1000

alpha = 0.1
theta_1, train_cost_history_1, valid_cost_history_1 =
gradient_descent(X_train, Y_train, X_test, Y_test, theta, alpha,
iterations) # 0.1 Learning Rate

alpha = 0.01
theta_2, train_cost_history_2, valid_cost_history_2 =
gradient_descent(X_train, Y_train, X_test, Y_test, theta, alpha,
iterations) # 0.01 Learning Rate

plt.plot(range(1, iterations+1), train_cost_history_1, color='blue',
label='Training Loss (0.1)')
plt.plot(range(1, iterations+1), valid_cost_history_1, color='red',
label='Validation Loss (0.1)')
plt.plot(range(1, iterations+1), train_cost_history_2, color='blue',
linestyle='dotted', label='Training Loss (0.01)')
plt.plot(range(1, iterations+1), valid_cost_history_2, color='red',
linestyle='dotted', label='Validation Loss (0.01)')
plt.rcParams["figure.figsize"] = (10,6)
plt.grid()
plt.xlabel('Number of iterations')
plt.ylabel('Cost (J)')
plt.title('Training vs Validation Loss')
plt.legend()
plt.show()

features = ['X_0', 'area', 'bedrooms', 'bathrooms', 'stories', 'parking']

print("\nBest Model Parameters (Theta) for Problem 2B - Normalized:
Learning Rate = 0.1\n")
theta_list = theta_1.flatten()


```

```

for label, value in zip(features, theta_list):
    print(f"{label:>12}: {value:.2f}")

```

### Figures / Results (2A):



### **Problem 2B**

#### Source Code (2B):

```

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import MinMaxScaler, StandardScaler
file_path = "Housing.csv"
sample = df = pd.read_csv(file_path)

inputs = ['area', 'bedrooms', 'bathrooms', 'stories', 'mainroad',
'guestroom', 'basement', 'hotwaterheating', 'airconditioning', 'parking',
'prefarea']

```

```

X = sample[inputs].values
Y = sample.values[:, 0]; #Price

binary_cols = ['mainroad', 'guestroom', 'basement', 'hotwaterheating',
'airconditioning', 'prefarea']
for col in binary_cols:
    sample[col] = sample[col].map({'yes': 1, 'no': 0})

m = len(Y)
X_0 = np.ones((m, 1))
X_cols = [sample['area'].values.reshape(m, 1),
           sample['bedrooms'].values.reshape(m, 1),
           sample['bathrooms'].values.reshape(m, 1),
           sample['stories'].values.reshape(m, 1),
           sample['mainroad'].values.reshape(m, 1),
           sample['guestroom'].values.reshape(m, 1),
           sample['basement'].values.reshape(m, 1),
           sample['hotwaterheating'].values.reshape(m, 1),
           sample['airconditioning'].values.reshape(m, 1),
           sample['parking'].values.reshape(m, 1),
           sample['prefarea'].values.reshape(m, 1)]

X = np.hstack([X_0] + X_cols)
theta = np.zeros(X.shape[1])

scaler = MinMaxScaler()
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.2,
random_state=42)
X_train[:, 1:] = scaler.fit_transform(X_train[:, 1:])
X_test[:, 1:] = scaler.transform(X_test[:, 1:])

# defining function for computing the cost for linear regression
def compute_cost(X, y, theta):
    m = len(y)
    predictions = X.dot(theta)
    errors = predictions - y
    sqrErrors = np.square(errors)
    J = 1 / (2 * m) * np.sum(sqrErrors)
    return J

```

```

# defining function for gradient descent algorithm
def gradient_descent(Xt, Yt, Xv, Yv, theta, alpha, iterations):
    train_cost_history = np.zeros(iterations)
    valid_cost_history = np.zeros(iterations)
    m = len(Yt)
    for i in range(iterations):
        predictions = Xt.dot(theta)
        errors = np.subtract(predictions, Yt)
        sum_delta = (alpha / m) * Xt.transpose().dot(errors)
        theta = theta - sum_delta
        train_cost_history[i] = compute_cost(Xt, Yt, theta)
        valid_cost_history[i] = compute_cost(Xv, Yv, theta)
    return theta, train_cost_history, valid_cost_history

theta = np.zeros(X.shape[1])
iterations = 1000

alpha = 0.1
theta_1, train_cost_history_1, valid_cost_history_1 =
gradient_descent(X_train, Y_train, X_test, Y_test, theta, alpha,
iterations) # 0.1 Learning Rate

alpha = 0.01
theta_2, train_cost_history_2, valid_cost_history_2 =
gradient_descent(X_train, Y_train, X_test, Y_test, theta, alpha,
iterations) # 0.01 Learning Rate

plt.plot(range(1, iterations+1), train_cost_history_1, color='blue',
label='Training Loss (0.1)')
plt.plot(range(1, iterations+1), valid_cost_history_1, color='red',
label='Validation Loss (0.1)')
plt.plot(range(1, iterations+1), train_cost_history_2, color='blue',
linestyle='dotted', label='Training Loss (0.01)')
plt.plot(range(1, iterations+1), valid_cost_history_2, color='red',
linestyle='dotted', label='Validation Loss (0.01)')
plt.rcParams["figure.figsize"] = (10, 6)
plt.grid()
plt.xlabel('Number of iterations')
plt.ylabel('Cost (J)')
plt.title('Training vs Validation Loss')

```

```


plt.legend()
plt.show()

features = ['X_0', 'area', 'bedrooms', 'bathrooms', 'stories', 'parking']

print("\nBest Model Parameters (Theta) for Problem 2A - Normalized:
Learning Rate = 0.1\n")
theta_list = theta_1.flatten()
for label, value in zip(features, theta_list):
    print(f"{label:>12}: {value:.2f}")

```

### Figures / Results (2B):



### **Problem 3A**

#### Source Code (3A):

```

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import MinMaxScaler, StandardScaler
file_path = "Housing.csv"
sample = df = pd.read_csv(file_path)

inputs = ['area', 'bedrooms', 'bathrooms', 'stories', 'parking'];
X = sample[inputs].values
Y = sample.values[:, 0]; #Price
m = len(Y)
X_0 = np.ones((m, 1))
X = np.hstack((X_0, X))
theta = np.zeros(X.shape[1])

scaler = MinMaxScaler()
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.2,
random_state=42)
X_train[:, 1:] = scaler.fit_transform(X_train[:, 1:])
X_test[:, 1:] = scaler.transform(X_test[:, 1:])

# defining function for computing the cost for linear regression
def compute_cost(X, y, theta):
    m = len(y)
    predictions = X.dot(theta)
    errors = predictions - y
    sqrErrors = np.square(errors)
    J = 1 / (2 * m) * np.sum(sqrErrors)
    return J

# defining function for gradient descent algorithm
def gradient_descent(Xt, Yt, Xv, Yv, theta, alpha, iterations, lam=0.1):
    train_cost_history = np.zeros(iterations)
    valid_cost_history = np.zeros(iterations)
    m = len(Yt)
    for i in range(iterations):
        predictions = Xt.dot(theta)
        errors = np.subtract(predictions, Yt)
        sum_delta = (alpha / m) * Xt.transpose().dot(errors)
        theta = theta - sum_delta
    return theta, train_cost_history, valid_cost_history

```

```

sum_delta[1:] += (lam/m) * theta[1:]
theta = theta - sum_delta
train_cost_history[i] = compute_cost(Xt, Yt, theta)
valid_cost_history[i] = compute_cost(Xv, Yv, theta)
return theta, train_cost_history, valid_cost_history

theta = np.zeros(X.shape[1])
iterations = 1000

alpha = 0.1
theta_1, train_cost_history_1, valid_cost_history_1 =
gradient_descent(X_train, Y_train, X_test, Y_test, theta, alpha,
iterations) # 0.1 Learning Rate

alpha = 0.01
theta_2, train_cost_history_2, valid_cost_history_2 =
gradient_descent(X_train, Y_train, X_test, Y_test, theta, alpha,
iterations) # 0.01 Learning Rate

plt.plot(range(1, iterations+1), train_cost_history_1, color='blue',
label='Training Loss (0.1)')
plt.plot(range(1, iterations+1), valid_cost_history_1, color='red',
label='Validation Loss (0.1)')
plt.plot(range(1, iterations+1), train_cost_history_2, color='blue',
linestyle='dotted', label='Training Loss (0.01)')
plt.plot(range(1, iterations+1), valid_cost_history_2, color='red',
linestyle='dotted', label='Validation Loss (0.01)')
plt.rcParams["figure.figsize"] = (10,6)
plt.grid()
plt.xlabel('Number of iterations')
plt.ylabel('Cost (J)')
plt.title('Training vs Validation Loss')
plt.legend()
plt.show()


features = ['X_0', 'area', 'bedrooms', 'bathrooms', 'stories', 'parking']

print("\nBest Model Parameters (Theta) for Problem 3B - Normalized +
Penalty: Learning Rate = 0.1\n")
theta_list = theta_1.flatten()

```

```
for label, value in zip(features, theta_list):
    print(f"{label:>12}: {value:.2f}")
```

### Figures / Results (3A):



### **Problem 3B**

#### Source Code (3B):

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import MinMaxScaler, StandardScaler
file_path = "Housing.csv"
sample = df = pd.read_csv(file_path)

inputs = ['area', 'bedrooms', 'bathrooms', 'stories', 'mainroad',
'guestroom', 'basement', 'hotwaterheating', 'airconditioning', 'parking',
'prefarea']
```

```

X = sample[inputs].values
Y = sample.values[:, 0]; #Price

binary_cols = ['mainroad', 'guestroom', 'basement', 'hotwaterheating',
'airconditioning', 'prefarea']
for col in binary_cols:
    sample[col] = sample[col].map({'yes': 1, 'no': 0})

m = len(Y)
X_0 = np.ones((m, 1))
X_cols = [sample['area'].values.reshape(m,1),
           sample['bedrooms'].values.reshape(m,1),
           sample['bathrooms'].values.reshape(m,1),
           sample['stories'].values.reshape(m,1),
           sample['mainroad'].values.reshape(m,1),
           sample['guestroom'].values.reshape(m,1),
           sample['basement'].values.reshape(m,1),
           sample['hotwaterheating'].values.reshape(m,1),
           sample['airconditioning'].values.reshape(m,1),
           sample['parking'].values.reshape(m,1),
           sample['prefarea'].values.reshape(m,1)]

X = np.hstack([X_0] + X_cols)
theta = np.zeros(X.shape[1])

scaler = MinMaxScaler()
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.2,
random_state=42)
X_train[:, 1:] = scaler.fit_transform(X_train[:, 1:])
X_test[:, 1:] = scaler.transform(X_test[:, 1:])

# defining function for computing the cost for linear regression
def compute_cost(X, y, theta):
    m = len(y)
    predictions = X.dot(theta)
    errors = predictions - y
    sqrErrors = np.square(errors)
    J = 1 / (2 * m) * np.sum(sqrErrors)
    return J

```

```

# defining function for gradient descent algorithm
def gradient_descent(Xt, Yt, Xv, Yv, theta, alpha, iterations, lam=0.1):
    train_cost_history = np.zeros(iterations)
    valid_cost_history = np.zeros(iterations)
    m = len(Yt)
    for i in range(iterations):
        predictions = Xt.dot(theta)
        errors = np.subtract(predictions, Yt)
        sum_delta = (alpha / m) * Xt.transpose().dot(errors)
        sum_delta[1:] += (lam/m) * theta[1:]
        theta = theta - sum_delta
        train_cost_history[i] = compute_cost(Xt, Yt, theta)
        valid_cost_history[i] = compute_cost(Xv, Yv, theta)
    return theta, train_cost_history, valid_cost_history

theta = np.zeros(X.shape[1])
iterations = 1000

alpha = 0.1
theta_1, train_cost_history_1, valid_cost_history_1 =
gradient_descent(X_train, Y_train, X_test, Y_test, theta, alpha,
iterations) # 0.1 Learning Rate

alpha = 0.01
theta_2, train_cost_history_2, valid_cost_history_2 =
gradient_descent(X_train, Y_train, X_test, Y_test, theta, alpha,
iterations) # 0.01 Learning Rate

plt.plot(range(1, iterations+1), train_cost_history_1, color='blue',
label='Training Loss (0.1)')
plt.plot(range(1, iterations+1), valid_cost_history_1, color='red',
label='Validation Loss (0.1)')
plt.plot(range(1, iterations+1), train_cost_history_2, color='blue',
linestyle='dotted', label='Training Loss (0.01)')
plt.plot(range(1, iterations+1), valid_cost_history_2, color='red',
linestyle='dotted', label='Validation Loss (0.01)')
plt.rcParams["figure.figsize"] = (10,6)
plt.grid()
plt.xlabel('Number of iterations')
plt.ylabel('Cost (J)')

```

```

plt.title('Training vs Validation Loss')
plt.legend()
plt.show()

features = ['X_0', 'area', 'bedrooms', 'bathrooms', 'stories', 'mainroad',
'guestroom', 'basement', 'hotwaterheating', 'airconditioning', 'parking',
'prefarea']

print("\nBest Model Parameters (Theta) for Problem 3B - Normalized +
Penalty: Learning Rate = 0.1\n")
theta_list = theta_1.flatten()
for label, value in zip(features, theta_list):
    print(f"{label:>12}: {value:.2f}")

```

Figure / Results (3B):

